

# ESCAPE ANALYSIS FOR MULTI-THREADED JAVALI

*Andrei Pârvu, Sebastian Wicki*

Department of Computer Science  
ETH Zürich  
Zürich, Switzerland

## ABSTRACT

Describe in concise words what you do, why you do it (not necessarily in this order), and the main result. The abstract has to be self-contained and readable for a person in the general area. You should write the abstract last.

## 1. INTRODUCTION

The basis of this project is an extended version of Javali with support for basic multi-threading, similar to Java's built-in `Thread` class. Additionally, we also added support for synchronization through locks. Having those additional language features also give more margin for potential optimizations, which we will discuss in the following paragraphs.

**Motivation.** Escape analysis can determine if an object reference escapes a certain dynamic scope, for example the scope of a thread or the scope of a stack frame. This information can be useful for a class of optimizations. In this project, we explored two optimizations: Stack allocation of objects and synchronization elision. Since Javali does not support modules or other forms of dynamic code loading, we can perform a whole-program escape analysis, typically yielding more precision.

In standard Javali, all non-primitive objects are allocated on the heap. The existing implementation even leaks those objects, as there is currently no garbage collection. Using escape analysis, the compiler can infer that certain objects will never outlive the current stack frame, thus those objects could be allocated on the stack. This has the advantage that the object will be cleaned up automatically once the function returns. Additionally, allocating data on the stack is typically faster than calling `malloc()` for heap allocation. This is especially true in multi-threaded programs, as the implementation of `malloc()` internally needs to synchronize as well, although modern implementations such as `TCMalloc` optimize this case heavily.

In this paper we provide an extended escape analysis for stack allocation that even allows stack allocation for objects that are shared with other threads, under the condition that the thread does not outlive the current function. This is often the case in the fork-join programming model, where a

function forks multiple child threads to work on in parallel on some data, and then waits for them to join.

As a second optimization, we implement synchronization removal on objects that never escape the current thread, i.e. that are never shared among multiple threads thus do not have to be locked in order to be accessed safely.

This optimization is useful for classes that are written to be thread-safe, i.e. use synchronization to ensure correctness, but are actually never used in a multi-threaded setting, which leads to unnecessary overhead. Library classes such as Java's `Vector` or `Hashtable` class are examples of thread-safe library classes that were later replaced with non-synchronized alternatives (namely `ArrayList` and `HashMap`) for better performance.

**Related work.** Next, you have to give a brief overview of related work. For a report like this, anywhere between 2 and 8 references. Briefly explain what they do. In the end contrast to what you do to make now precisely clear what your contribution is.

## 2. BACKGROUND

In this section we briefly present the programming interface and semantics of our threading extension. In the second part, we formally introduce escape analysis.

**Threading and Synchronization in Javali.** The interface for multi-threading in Javali is heavily inspired by Java. In order to create a new thread, the programmer has to create a new class that inherits from the built-in `Thread` class and overwrite the `run()` method to provide a thread entry point. The `Thread` class contains the following methods:

- run()** Thread entry point. To be overwritten by clients.
- start()** Starts a new thread. Like in Java it is an error to call this method more than once on the same receiver.
- join()** Blocks the caller until the thread exits. It is an error to call this on a thread that was never started or was already joined.

Similar to Java, our extension to Javali has a built-in lock and conditional variable within every object. A notable difference to Java however is that there is no special syntax for locking, the built-in `Object` just provides the following methods:

**lock()** Locks the receiver. If the object was already locked by another thread, the current thread is blocked until the object is unlocked. The lock is re-entrant in the sense that a thread can call this method on objects on which it already own the lock.

**unlock()** Unlocks the object. Also re-entrant in the sense that in order to fully unlock an object, the lock owner needs to call `unlock()` as many times as `lock()` was called. It is an error to call `unlock()` without owning the lock.

**wait()** Causes the calling thread to wait until another thread calls `notify()` on the object.

**notify()** Wakes up another thread waiting on this object.

#### Escape Analysis.

Give a short, self-contained summary of necessary background information. For example, assume you present an implementation of sorting algorithms. You could organize into sorting definition, algorithms considered, and asymptotic runtime statements. The goal of the background section is to make the paper self-contained for an audience as large as possible. As in every section you start with a very brief overview of the section. Here it could be as follows: In this section we formally define the sorting problem we consider and introduce the algorithms we use including a cost analysis.

**Sorting.** Precisely define sorting problem you consider.

**Sorting algorithms.** Explain the algorithm you use including their costs.

As an aside, don't talk about "the complexity of the algorithm." It's incorrect, problems have a complexity, not algorithms.

### 3. YOUR PROPOSED METHOD

Now comes the "beef" of the report, where you explain what you did. Again, organize it in paragraphs with titles. As in every section you start with a very brief overview of the section.

In this section, structure is very important so one can follow the technical content.

Mention and cite any external resources that you used including libraries or other code.




Fig. 1. Running times matrix multiplication

## 4. EXPERIMENTAL RESULTS

**Experimental setup.** For the experiments we used two different machines. The first one, running Ubuntu 14.04 with a i5-3317U, 1.70GHz processor and 4Gb RAM.

We tried to analyze three different aspects of our project: firstly, the running time improvement of just allocating the unescaped objects on the stack. Secondly, we wanted to see the time benefit of working with stack-allocated objects instead of heap-allocated ones. And thirdly, we tried to observe the how lock-removal affects single-threaded programs.


**Results.** For the first experiment, we just allocated various objects without using them afterwards, in order to see how the running time differs for both cases. TODO: run the actual experiment :)

In the second experiment, we wrote a basic matrix multiplication program with the intent of observing how execution time changes when using stack allocation. The results can be seen in Figure 1. It can be observed that the stack allocation improves running time with almost 20%, with the obvious draw-back that more stack memory is used.


For the third experiment, we implemented a hash map, which simulates string hashing (because Javali doesn't have strings we used ASCII arrays). The hash map is implemented to be thread-safe, so we ran the experiment with one thread and lock removal. Results can be seen in Figure 2; with lock removal, running time is approximately 10% faster than the unanalyzed version.

#### Comments:

- Create very readable, attractive plots (do 1 column, not 2 column plots for this report) with readable font size. However, the font size should also not be too large; typically it is smaller than the text font size. An example is in Fig. 3 (of course you can have a different style).



**Fig. 2.** Running times hash map



**Fig. 3.** Performance of four single precision implementations of the discrete Fourier transform. The operations count is roughly the same. The labels in this plot are maybe a little bit too small.

- Every plot answers a question. You state this question and extract the answer from the plot in its discussion.
- Every plot should be referenced and discussed.

## 5. CONCLUSIONS

Here you need to summarize what you did and why this is important. *Do not take the abstract* and put it in the past tense. Remember, now the reader has (hopefully) read the report, so it is a very different situation from the abstract. Try to highlight important results and say the things you really want to get across such as high-level statements (e.g., we believe that .... is the right approach to .... Even though we only considered x, the .... technique should be applicable ....) You can also formulate next steps if you want. Be brief. After the conclusions there are only the references.

## 6. FURTHER COMMENTS

Here we provide some further tips.

### Further general guidelines.

- For short papers, to save space, I use paragraph titles instead of subsections, as shown in the introduction.
- It is generally a good idea to break sections into such smaller units for readability and since it helps you to (visually) structure the story.
- The above section titles should be adapted to more precisely reflect what you do.
- Each section should be started with a very short summary of what the reader can expect in this section. Nothing more awkward as when the story starts and one does not know what the direction is or the goal.
- Make sure you define every acronym you use, no matter how convinced you are the reader knows it.
- Always spell-check before you submit (to us in this case).
- Be picky. When writing a paper you should always strive for very high quality. Many people may read it and the quality makes a big difference. In this class, the quality is part of the grade.
- Books helping you to write better: [?] and [?].
- Conversion to pdf (latex users only):  
`dvips -o conference.ps -t letter -Ppdf -G0 conference.dvi`  
and then  
`ps2pdf conference.ps`

**Graphics.** For plots that are not images *never* generate the bitmap formats jpeg, gif, bmp, tif. Use eps, which means encapsulate postscript. It is scalable since it is a vector graphic description of your graph. E.g., from Matlab, you can export to eps.

The format pdf is also fine for plots (you need pdf<sub>l</sub>atex then), but only if the plot was never before in the format jpeg, gif, bmp, tif.

## 7. REFERENCES

- [1] J. Whaley, M. Rinard, *Compositional Pointer and Escape Analysis for Java Programs* 1999
- [2] E. Ruf, *Effective Synchronization Removal for Java* 2000
- [3] V.P. Ranganath, J. Hatcliff *Pruning interference and ready dependences for slicing concurrent Java programs* 2004